

```

# Install necessary dependencies
!pip install -q opencv-python-headless
!pip install -q matplotlib
!pip install -q numpy
!pip install -q scikit-image
!pip install -q tensorflow
!pip install -q albumentations

import os
import cv2
import numpy as np
import matplotlib.pyplot as plt
from google.colab import files
from sklearn.cluster import KMeans
from skimage import exposure, filters, color
import random
from typing import List, Tuple, Dict, Any, Optional
import glob
import time

# Set random seeds for reproducibility
np.random.seed(42)
random.seed(42)

class CarDamagePreprocessor:
    """
    A class for preprocessing damaged car images to prepare them for damage detection
    and classification models.
    """

    def __init__(self,
                 target_size: Tuple[int, int] = (512, 512),
                 normalize: bool = True,
                 clahe_clip_limit: float = 2.0,
                 clahe_grid_size: Tuple[int, int] = (8, 8)):
        """
        Initialize the damaged car image preprocessor.

        Args:
            target_size: Output size for processed images (height, width)
            normalize: Whether to normalize pixel values to [0,1]
            clahe_clip_limit: Clip limit for CLAHE contrast enhancement
            clahe_grid_size: Grid size for CLAHE contrast enhancement
        """
        self.target_size = target_size
        self.normalize = normalize
        self.clahe = cv2.createCLAHE(clipLimit=clahe_clip_limit,
                                     tileGridSize=clahe_grid_size)

    def load_image(self, image_path: str) -> np.ndarray:
        """
        Load an image from a file path.

        Args:
            image_path: Path to the image file

        Returns:
            The loaded image as a numpy array
        """
        image = cv2.imread(image_path)
        if image is None:
            raise ValueError(f"Failed to load image from {image_path}")
        return cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

    def standardize_image(self, image: np.ndarray) -> np.ndarray:
        """
        Resize and standardize an image.

        Args:
            image: Input image as numpy array

        Returns:
            Standardized image
        """
        # Resize to target size
        resized = cv2.resize(image, (self.target_size[1], self.target_size[0]))

        # Normalize pixel values if requested
        if self.normalize:

```

```

        return resized.astype(np.float32) / 255.0
    return resized

def remove_background(self, image: np.ndarray,
                      threshold: int = 25,
                      blur_size: int = 5) -> Tuple[np.ndarray, np.ndarray]:
    """
    Remove the background from a car image to focus on the vehicle.
    Uses GrabCut algorithm for automatic foreground extraction.

    Args:
        image: Input image as numpy array
        threshold: Threshold for background removal
        blur_size: Size of the blur kernel for preprocessing

    Returns:
        Tuple of (processed image with background removed, mask)
    """
    # Create a copy of the image
    img = image.copy()

    # Convert to RGB if needed
    if len(img.shape) == 2:
        img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)

    # Ensure image is uint8 for GrabCut (required by OpenCV)
    if img.dtype == np.float32:
        img = (img * 255).astype(np.uint8)
    elif img.dtype != np.uint8:
        img = img.astype(np.uint8)

    # Initial mask creation
    mask = np.zeros(img.shape[:2], np.uint8)

    # Background and foreground models
    bgd_model = np.zeros((1, 65), np.float64)
    fgd_model = np.zeros((1, 65), np.float64)

    # Define rough ROI around the image center assuming car is in the middle
    margin = 50
    rect = (margin, margin, img.shape[1]-2*margin, img.shape[0]-2*margin)

    try:
        # Apply GrabCut
        cv2.grabCut(img, mask, rect, bgd_model, fgd_model, 5, cv2.GC_INIT_WITH_RECT)

        # Convert mask
        mask2 = np.where((mask==2) | (mask==0), 0, 1).astype('uint8')
    except cv2.error:
        # Fallback if GrabCut fails
        print("GrabCut failed. Using basic thresholding as fallback.")
        gray = cv2.cvtColor(img, cv2.COLOR_RGB2GRAY)
        blurred = cv2.GaussianBlur(gray, (blur_size, blur_size), 0)
        _, mask2 = cv2.threshold(blurred, threshold, 1, cv2.THRESH_BINARY)

    # Apply the mask to the image
    result = img * mask2[:, :, np.newaxis]

    # Convert back to original format if needed
    if image.dtype == np.float32:
        result = result.astype(np.float32) / 255.0

    return result, mask2

def detect_roi(self, image: np.ndarray,
               mask: Optional[np.ndarray] = None) -> Tuple[np.ndarray, Tuple[int, int, int, int]]:
    """
    Detect the region of interest (ROI) containing the damaged car.

    Args:
        image: Input image
        mask: Optional mask from background removal

    Returns:
        Tuple of (cropped image containing ROI, bounding box coordinates)
    """
    # If mask is provided, use it to find contours
    if mask is not None:
        contours, _ = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
        if contours:
            # Find the largest contour (assumed to be the car)
            largest_contour = max(contours, key=cv2.contourArea)
            x, y, w, h = cv2.boundingRect(largest_contour)

```

```

    # Add some padding
    padding = 10
    x = max(0, x - padding)
    y = max(0, y - padding)
    w = min(image.shape[1] - x, w + 2*padding)
    h = min(image.shape[0] - y, h + 2*padding)

    # Crop the image to the bounding box
    cropped = image[y:y+h, x:x+w]
    return cropped, (x, y, w, h)

# If no mask or no contours found, use edge detection as fallback
gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY) if len(image.shape) > 2 else image
blurred = cv2.GaussianBlur(gray, (5, 5), 0)
edges = cv2.Canny(blurred, 50, 150)

# Find contours in the edge map
contours, _ = cv2.findContours(edges, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)

if contours:
    # Combine all contours to find the overall bounding box
    all_points = np.concatenate([cnt for cnt in contours])
    x, y, w, h = cv2.boundingRect(all_points)

    # Add some padding
    padding = 20
    x = max(0, x - padding)
    y = max(0, y - padding)
    w = min(image.shape[1] - x, w + 2*padding)
    h = min(image.shape[0] - y, h + 2*padding)

    # Crop the image to the bounding box
    cropped = image[y:y+h, x:x+w]
    return cropped, (x, y, w, h)

# If all else fails, return the original image
return image, (0, 0, image.shape[1], image.shape[0])

def reduce_noise(self, image: np.ndarray,
                 method: str = 'gaussian',
                 kernel_size: int = 5) -> np.ndarray:
    """
    Apply noise reduction to an image.

    Args:
        image: Input image
        method: Noise reduction method ('gaussian', 'median', 'bilateral')
        kernel_size: Size of the kernel for noise reduction

    Returns:
        Noise-reduced image
    """
    if method == 'gaussian':
        return cv2.GaussianBlur(image, (kernel_size, kernel_size), 0)
    elif method == 'median':
        return cv2.medianBlur(image, kernel_size)
    elif method == 'bilateral':
        if len(image.shape) > 2 and image.dtype == np.float32:
            # Convert to 8-bit for bilateral filter
            temp = (image * 255).astype(np.uint8)
            result = cv2.bilateralFilter(temp, kernel_size, 75, 75)
            return result.astype(np.float32) / 255.0
        else:
            return cv2.bilateralFilter(image, kernel_size, 75, 75)
    else:
        raise ValueError(f"Unknown noise reduction method: {method}")

def enhance_contrast(self, image: np.ndarray,
                    method: str = 'clahe') -> np.ndarray:
    """
    Enhance contrast in an image to make damage more visible.

    Args:
        image: Input image
        method: Contrast enhancement method ('clahe', 'histeq', 'adapthist')

    Returns:
        Contrast-enhanced image
    """
    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:

```

```

        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Scale to 0-255 if normalized
    if gray.dtype == np.float32:
        gray = (gray * 255).astype(np.uint8)

    if method == 'clahe':
        enhanced = self.clahe.apply(gray)
    elif method == 'histeq':
        enhanced = cv2.equalizeHist(gray)
    elif method == 'adapthist':
        enhanced = exposure.equalize_adapthist(gray, clip_limit=0.03)
        enhanced = (enhanced * 255).astype(np.uint8)
    else:
        raise ValueError(f"Unknown contrast enhancement method: {method}")

    # If input was RGB, convert back to RGB
    if len(image.shape) > 2:
        # Create a 3-channel image where each channel has the enhanced data
        enhanced_rgb = np.zeros_like(image)
        if image.dtype == np.float32:
            enhanced_rgb[:, :, 0] = enhanced.astype(np.float32) / 255.0
            enhanced_rgb[:, :, 1] = enhanced.astype(np.float32) / 255.0
            enhanced_rgb[:, :, 2] = enhanced.astype(np.float32) / 255.0
        else:
            enhanced_rgb[:, :, 0] = enhanced
            enhanced_rgb[:, :, 1] = enhanced
            enhanced_rgb[:, :, 2] = enhanced
        return enhanced_rgb

    # Return the enhanced grayscale image
    if image.dtype == np.float32:
        return enhanced.astype(np.float32) / 255.0
    return enhanced

def detect_edges(self, image: np.ndarray,
                 method: str = 'canny',
                 low_threshold: int = 50,
                 high_threshold: int = 150) -> np.ndarray:
    """
    Detect edges in an image to highlight damage areas.

    Args:
        image: Input image
        method: Edge detection method ('canny', 'sobel', 'scharr')
        low_threshold: Low threshold for Canny edge detection
        high_threshold: High threshold for Canny edge detection

    Returns:
        Edge map
    """
    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:
        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Scale to 0-255 if normalized
    if gray.dtype == np.float32:
        gray = (gray * 255).astype(np.uint8)

    # Apply Gaussian blur to reduce noise
    blurred = cv2.GaussianBlur(gray, (5, 5), 0)

    if method == 'canny':
        edges = cv2.Canny(blurred, low_threshold, high_threshold)
    elif method == 'sobel':
        sobelx = cv2.Sobel(blurred, cv2.CV_64F, 1, 0, ksize=3)
        sobely = cv2.Sobel(blurred, cv2.CV_64F, 0, 1, ksize=3)
        edges = np.sqrt(sobelx**2 + sobely**2)
        edges = cv2.normalize(edges, None, 0, 255, cv2.NORM_MINMAX, cv2.CV_8U)
    elif method == 'scharr':
        scharrx = cv2.Scharr(blurred, cv2.CV_64F, 1, 0)
        scharry = cv2.Scharr(blurred, cv2.CV_64F, 0, 1)
        edges = np.sqrt(scharrx**2 + scharry**2)
        edges = cv2.normalize(edges, None, 0, 255, cv2.NORM_MINMAX, cv2.CV_8U)
    else:
        raise ValueError(f"Unknown edge detection method: {method}")

    # Return the edge map
    if image.dtype == np.float32:

```

```

    return edges.astype(np.float32) / 255.0
    return edges

def segment_damage(self, image: np.ndarray,
                    edge_map: np.ndarray = None,
                    threshold: float = 0.3) -> np.ndarray:
    """
    Simple damage segmentation based on edge information.
    This is a basic approach that can be refined with ML techniques.

    Args:
        image: Input image
        edge_map: Edge map from edge detection
        threshold: Threshold for damage segmentation

    Returns:
        Mask highlighting potential damage areas
    """
    if edge_map is None:
        edge_map = self.detect_edges(image)

    # Threshold the edge map to get binary mask
    if edge_map.dtype == np.float32:
        mask = (edge_map > threshold).astype(np.uint8)
    else:
        mask = (edge_map > threshold * 255).astype(np.uint8)

    # Apply morphological operations to clean up the mask
    kernel = np.ones((5, 5), np.uint8)
    mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, kernel)
    mask = cv2.morphologyEx(mask, cv2.MORPH_OPEN, kernel)

    # Label connected components
    num_labels, labels = cv2.connectedComponents(mask)

    # Filter out small regions
    min_size = 50
    for i in range(1, num_labels):
        if np.sum(labels == i) < min_size:
            mask[labels == i] = 0

    return mask

def extract_features(self, image: np.ndarray,
                     mask: Optional[np.ndarray] = None) -> Dict[str, Any]:
    """
    Extract features from the image for damage analysis.

    Args:
        image: Input image
        mask: Optional mask to focus on specific regions

    Returns:
        Dictionary of extracted features
    """
    # Convert to grayscale if image is RGB
    if len(image.shape) > 2:
        gray = cv2.cvtColor(image, cv2.COLOR_RGB2GRAY)
    else:
        gray = image.copy()

    # Apply mask if provided
    if mask is not None:
        masked_gray = cv2.bitwise_and(gray, gray, mask=mask)
    else:
        masked_gray = gray

    # Scale to 0-255 if normalized
    if masked_gray.dtype == np.float32:
        masked_gray = (masked_gray * 255).astype(np.uint8)

    # Extract features
    features = {}

    # Basic statistics
    if np.any(masked_gray > 0):
        features['mean'] = np.mean(masked_gray[masked_gray > 0])
        features['std'] = np.std(masked_gray[masked_gray > 0])
        features['min'] = np.min(masked_gray[masked_gray > 0])
        features['max'] = np.max(masked_gray[masked_gray > 0])
    else:
        features['mean'] = 0

```

```

features['std'] = 0
features['min'] = 0
features['max'] = 0

# Histogram
hist = cv2.calcHist([masked_gray], [0], None, [256], [0, 256])
features['histogram'] = hist.flatten()

# Texture features using Haralick texture features (calculated manually)
if np.any(masked_gray > 0):
    # Convert to uint8 for texture analysis
    masked_gray_uint8 = masked_gray.astype(np.uint8)

    # Calculate gradient magnitude as a simple texture feature
    sobelx = cv2.Sobel(masked_gray_uint8, cv2.CV_64F, 1, 0, ksize=3)
    sobely = cv2.Sobel(masked_gray_uint8, cv2.CV_64F, 0, 1, ksize=3)
    gradient_magnitude = np.sqrt(sobelx**2 + sobely**2)

    features['gradient_mean'] = np.mean(gradient_magnitude)
    features['gradient_std'] = np.std(gradient_magnitude)

    # Calculate local binary pattern (simple version)
    def local_binary_pattern(image, points=8, radius=1):
        rows, cols = image.shape
        result = np.zeros((rows-2*radius, cols-2*radius), dtype=np.uint8)
        for i in range(radius, rows-radius):
            for j in range(radius, cols-radius):
                center = image[i, j]
                pattern = 0
                for p in range(points):
                    angle = 2 * np.pi * p / points
                    x = j + int(round(radius * np.cos(angle)))
                    y = i + int(round(radius * np.sin(angle)))
                    if image[y, x] >= center:
                        pattern |= (1 << p)
                result[i-radius, j-radius] = pattern
        return result

    try:
        # Only compute LBP on a smaller region if image is large
        if masked_gray_uint8.shape[0] > 100 and masked_gray_uint8.shape[1] > 100:
            center_y, center_x = masked_gray_uint8.shape[0] // 2, masked_gray_uint8.shape[1] // 2
            roi_size = 50
            roi = masked_gray_uint8[
                max(0, center_y - roi_size):min(masked_gray_uint8.shape[0], center_y + roi_size),
                max(0, center_x - roi_size):min(masked_gray_uint8.shape[1], center_x + roi_size)
            ]
            lbp = local_binary_pattern(roi)
        else:
            lbp = local_binary_pattern(masked_gray_uint8)

        lbp_hist = cv2.calcHist([lbp], [0], None, [256], [0, 256])
        features['lbp_histogram'] = lbp_hist.flatten()
        features['lbp_entropy'] = -np.sum((lbp_hist / np.sum(lbp_hist)) *
                                           np.log2(lbp_hist / np.sum(lbp_hist) + 1e-10))

    except Exception as e:
        print(f"LBP calculation error: {e}")
        features['lbp_histogram'] = np.zeros(256)
        features['lbp_entropy'] = 0

else:
    features['gradient_mean'] = 0
    features['gradient_std'] = 0
    features['lbp_histogram'] = np.zeros(256)
    features['lbp_entropy'] = 0

# SIFT features (keypoints)
try:
    if np.any(masked_gray > 0):
        sift = cv2.SIFT_create()
        keypoints, descriptors = sift.detectAndCompute(masked_gray, None)
        features['num_keypoints'] = len(keypoints)
        features['keypoints'] = keypoints
        features['descriptors'] = descriptors if descriptors is not None else np.array([])
    else:
        features['num_keypoints'] = 0
        features['keypoints'] = []
        features['descriptors'] = np.array([])
except Exception as e:
    print(f"SIFT feature extraction error: {e}")
    features['num_keypoints'] = 0
    features['keypoints'] = []
    features['descriptors'] = np.array([])

```

```
return features
```

```
def augment_data(self, image: np.ndarray,
                  num_augmentations: int = 5) -> List[np.ndarray]:
    """
    Generate augmented versions of the input image for training.

    Args:
        image: Input image
        num_augmentations: Number of augmented images to generate

    Returns:
        List of augmented images
    """
    augmented_images = []

    # Define some augmentation functions
    def random_brightness_contrast(img, brightness_range=(-0.2, 0.2), contrast_range=(-0.2, 0.2)):
        # Brightness adjustment
        brightness = np.random.uniform(brightness_range[0], brightness_range[1])
        adjusted = img.astype(np.float32) + brightness

        # Contrast adjustment
        contrast = np.random.uniform(contrast_range[0], contrast_range[1]) + 1.0
        adjusted = adjusted * contrast

        # Clip values to valid range
        adjusted = np.clip(adjusted, 0, 1.0 if img.dtype == np.float32 else 255)
        return adjusted.astype(img.dtype)

    def random_noise(img, var=0.01):
        # Add Gaussian noise
        if img.dtype == np.float32:
            noise = np.random.normal(0, var**0.5, img.shape)
            noisy = img + noise
            return np.clip(noisy, 0, 1.0).astype(np.float32)
        else:
            noise = np.random.normal(0, var**0.5 * 255, img.shape).astype(np.int16)
            noisy = img.astype(np.int16) + noise
            return np.clip(noisy, 0, 255).astype(np.uint8)

    def random_rotation(img, angle_range=(-15, 15)):
        # Random rotation
        angle = np.random.uniform(angle_range[0], angle_range[1])
        rows, cols = img.shape[:2]
        M = cv2.getRotationMatrix2D((cols/2, rows/2), angle, 1)
        return cv2.warpAffine(img, M, (cols, rows))

    def random_flip(img):
        # Random horizontal flip
        if np.random.random() > 0.5:
            return cv2.flip(img, 1)
        return img

    def random_crop(img, crop_factor_range=(0.8, 0.95)):
        factor = np.random.uniform(crop_factor_range[0], crop_factor_range[1])
        h, w = img.shape[:2]
        crop_h, crop_w = int(h * factor), int(w * factor)
        start_h = np.random.randint(0, h - crop_h + 1)
        start_w = np.random.randint(0, w - crop_w + 1)
        cropped = img[start_h:start_h+crop_h, start_w:start_w+crop_w]
        return cv2.resize(cropped, (w, h))

    # Define augmentation pipeline with probabilities
    augmentation_functions = [
        (random_brightness_contrast, 0.7),
        (random_noise, 0.5),
        (random_rotation, 0.5),
        (random_flip, 0.5),
        (random_crop, 0.5)
    ]

    for _ in range(num_augmentations):
        # Start with a copy of the original image
        augmented = image.copy()

        # Apply random augmentations based on probability
        for aug_func, prob in augmentation_functions:
            if np.random.random() < prob:
                augmented = aug_func(augmented)

        augmented_images.append(augmented)
```

```

    return augmented_images

def visualize_preprocessing(self, original: np.ndarray,
                           processed_results: Dict[str, np.ndarray]) -> None:
    """
    Visualize the preprocessing steps.

    Args:
        original: Original image
        processed_results: Dictionary of processed images
    """
    # Determine number of steps
    n_steps = len(processed_results) + 1 # +1 for original

    # Create figure with subplots
    fig, axes = plt.subplots(1, n_steps, figsize=(20, 5))

    # Plot original image
    axes[0].imshow(original)
    axes[0].set_title('Original')
    axes[0].axis('off')

    # Plot processed results
    for i, (title, img) in enumerate(processed_results.items(), 1):
        # Handle different image types
        if len(img.shape) == 2: # Grayscale or mask
            if img.dtype == bool:
                img = img.astype(np.uint8) * 255

            # Display as grayscale
            axes[i].imshow(img, cmap='gray')
        else:
            # Display as RGB
            if img.dtype == np.float32 and np.max(img) <= 1.0:
                axes[i].imshow(img)
            else:
                axes[i].imshow(img.astype(np.uint8))

        axes[i].set_title(title)
        axes[i].axis('off')

    plt.tight_layout()
    plt.show()

def process_image(self, image: np.ndarray,
                  visualize: bool = False) -> Dict[str, Any]:
    """
    Process a single image through the entire pipeline.

    Args:
        image: Input image
        visualize: Whether to visualize the preprocessing steps

    Returns:
        Dictionary of processed images and features
    """
    results = {}

    # Standardize image
    std_image = self.standardize_image(image)
    results['standardized'] = std_image

    # Remove background
    bg_removed, mask = self.remove_background(std_image)
    results['background_removed'] = bg_removed
    results['background_mask'] = mask

    # Detect ROI
    roi, bbox = self.detect_roi(bg_removed, mask)
    results['roi'] = roi
    results['bbox'] = bbox

    # Reduce noise
    denoised = self.reduce_noise(roi, method='bilateral')
    results['denoised'] = denoised

    # Enhance contrast
    enhanced = self.enhance_contrast(denoised)
    results['enhanced'] = enhanced

    # Detect edges
    edges = self.detect_edges(enhanced)

```



```

results['edges'] = edges

# Segment damage
damage_mask = self.segment_damage(enhanced, edges)
results['damage_mask'] = damage_mask

# Extract features
features = self.extract_features(enhanced, damage_mask)
results['features'] = features

# Visualize if requested
if visualize:
    vis_results = {
        'Background Removed': bg_removed,
        'ROI': roi,
        'Denoised': denoised,
        'Enhanced': enhanced,
        'Edges': edges,
        'Damage Mask': damage_mask
    }
    self.visualize_preprocessing(image, vis_results)

return results

def process_directory(self, directory_path: str,
                      output_dir: str = None,
                      visualize: bool = False) -> Dict[str, Dict[str, Any]]:
    """
    Process all images in a directory.

    Args:
        directory_path: Path to directory containing images
        output_dir: Path to directory to save processed images
        visualize: Whether to visualize the preprocessing steps

    Returns:
        Dictionary mapping image filenames to processing results
    """
    # Create output directory if specified
    if output_dir is not None:
        os.makedirs(output_dir, exist_ok=True)

    results = {}

    # Get all image files
    image_files = []
    for ext in ['*.jpg', '*.jpeg', '*.png', '*.bmp']:
        image_files.extend(glob.glob(os.path.join(directory_path, ext)))
        image_files.extend(glob.glob(os.path.join(directory_path, ext.upper())))

    print(f"Found {len(image_files)} images in {directory_path}")

    # Process each image
    for image_file in image_files:
        try:
            # Load image
            image = self.load_image(image_file)

            # Process image
            result = self.process_image(image, visualize=visualize)

            # Save processed images if output directory is specified
            if output_dir is not None:
                # Get base filename without extension
                basename = os.path.splitext(os.path.basename(image_file))[0]

                # Save each processed image
                for name, img in result.items():
                    if isinstance(img, np.ndarray):
                        # Create image file path
                        img_path = os.path.join(output_dir, f"{basename}_{name}.png")

                        # Convert to uint8 if needed
                        if img.dtype == np.float32:
                            img = (img * 255).astype(np.uint8)

                        # Save the image
                        if len(img.shape) == 2:
                            cv2.imwrite(img_path, img)
                        else:
                            cv2.imwrite(img_path, cv2.cvtColor(img, cv2.COLOR_RGB2BGR))

```

```

# Store results
results[os.path.basename(image_file)] = result

except Exception as e:
    print(f"Error processing {image_file}: {e}")

return results

# Example usage
def main():
    """
    Example usage of the CarDamagePreprocessor.
    """
    # Create preprocessor
    preprocessor = CarDamagePreprocessor()

    # Check if images already exist in the environment
    import os
    existing_images = [f for f in os.listdir() if f.lower().endswith(('.png', '.jpg', '.jpeg', '.bmp'))]

    if existing_images:
        print(f"Found {len(existing_images)} images in the current directory.")
        image_files = existing_images
    else:
        # Allow user to upload images
        print("Please upload one or more damaged car images.")
        uploaded = files.upload()
        image_files = list(uploaded.keys())

    # Process images
    for filename in image_files:
        try:
            print(f"Processing {filename}...")

            # Load image
            image = cv2.imread(filename)
            if image is None:
                print(f"Error: Could not read image {filename}")
                continue

            image = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

            # Process image with error handling
            try:
                result = preprocessor.process_image(image, visualize=True)

                # Display features
                print(f"Image features:")
                for key, value in result['features'].items():
                    if key in ['histogram', 'keypoints', 'descriptors']:
                        if isinstance(value, np.ndarray):
                            print(f"  {key}: [array with shape {value.shape}]")
                        else:
                            print(f"  {key}: [array with {len(value)} elements]")
                    else:
                        print(f"  {key}: {value}")

                print("\n")

                # Demonstrate augmentation with the first successful image
                print("Generating data augmentations...")
                augmented_images = preprocessor.augment_data(image, num_augmentations=5)

                # Display augmented images
                plt.figure(figsize=(15, 10))
                plt.subplot(2, 3, 1)
                plt.imshow(image)
                plt.title("Original")
                plt.axis('off')

                for i, aug_img in enumerate(augmented_images, 1):
                    plt.subplot(2, 3, i+1)
                    plt.imshow(aug_img)
                    plt.title(f"Augmentation {i}")
                    plt.axis('off')

                plt.tight_layout()
                plt.show()

                # Only process one image for demonstration
                break

            except Exception as e:

```

```
print(f"Error during image processing: {str(e)}")
import traceback
traceback.print_exc()

except Exception as e:
    print(f"Error with image {filename}: {str(e)}")
    import traceback
    traceback.print_exc()

if __name__ == "__main__":
    main()
```

Saving auto-3734396_1280.jpg to auto-3734396_1280.jpg
Processing auto-3734396_1280.jpg...

